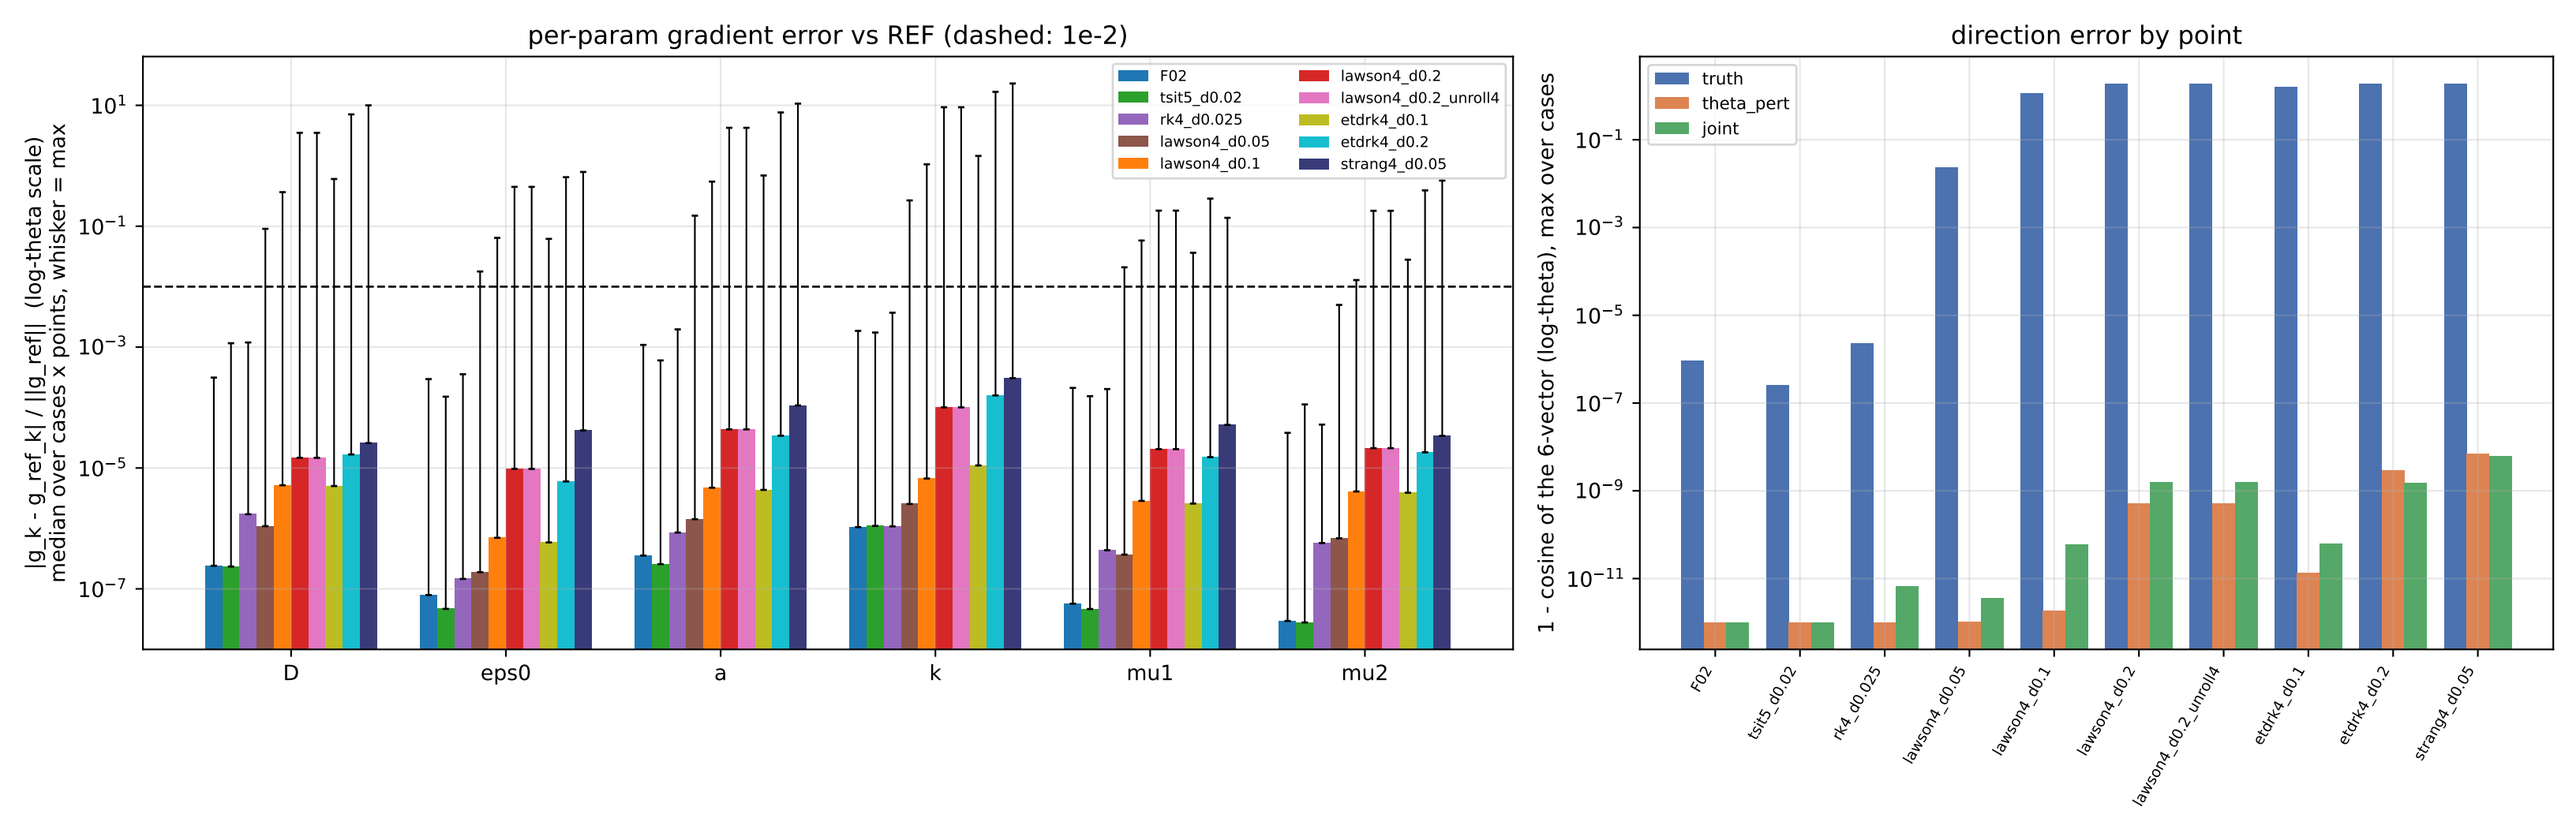
per-param gradient error vs REF (dashed: 1e-2)
|g_k - g_ref_k| / ||g_ref|| (log-theta scale)
median over cases x points, whisker = max
F02
tsit5_d0.02
rk4_d0.025
lawson4_d0.05
lawson4_d0.1
lawson4_d0.2
lawson4_d0.2_unroll4
etdrk4_d0.1
etdrk4_d0.2
strang4_d0.05
D
eps0
a
k
mu1
mu2
direction error by point
1 - cosine of the 6-vector (log-theta), max over cases
truth
theta_pert
joint